

SAPIENZA UNIVERSITÀ DI
ROMA



Machine Learning

Homework 2

Ivan MAZZANTI
2122298

1 Preliminary remarks

To properly solve this problem, it is very important to make some initial considerations before moving on to the architecture design of the convolutional neural network.

1.1 Unbalanced and limited dataset

Let us examine the training set.

- **Class 0** = 1000 images;
- **Class 1** = 1500 images;
- **Class 2** = 1500 images;
- **Class 3** = 2000 images;
- **Class 4** = 369 images;

As we can see, the training set contains a total of 6369 images. The images associated with class 3 are the most numerous, with a total of 2000. Those associated with class 1 have exactly half of the elements of class 3 and, finally, we have only 369 images associated with class 4. We can observe that the training set is limited and unbalanced. Moreover, this situation is likely to get worse, as part of the training set will have to be reserved to become the validation set. Having a limited training set and with underrepresented classes increases the difficulty of obtaining good results, and if we consider that using a CNN for this problem will not give us optimal performance as it would be better to use Reinforcement Learning (because the history of events is also important in this problem), the goal becomes even more complicated. To improve performance, to correctly train the model, and to prevent overfitting, it will be very important to consider dataset augmentation techniques, not create a too complex CNN, and evaluate regularization techniques. Finally, when evaluating the model, it will be important to consider not only the accuracy but also the other fundamental classification evaluation metrics.

1.2 Chosen approaches and hyperparameters

To solve this task, I decided to create two different CNN architectures and to use two different optimizers: Adam and SGD with momentum. The tests performed and the results obtained will now be described. The library I used for the implementation is PyTorch.

2 First approach

2.1 Data augmentation

As a first step, given the considerations made previously, I performed data augmentation by applying transformations to the images of the training set. These transformations include: **random horizontal flipping, random rotation of 5 degrees and random resize crop (similar to a zoom, but it randomly crops a part of the image and resizes it as specified)**. I decided not to use too "aggressive" transformations or to add any others because doing so would cause overfitting. I have therefore reserved 1/3 of the training set for the validation set and applied data augmentation only on the training set. The snippet of code that performs the transformations is the following:

```
transform = Compose([
    RandomHorizontalFlip(),
    RandomRotation(5),
    transforms.RandomResizedCrop(96,
                                  scale = (0.8, 1.0)),
    transforms.ToTensor()
])
```

2.2 CNN architecture design

For the first architecture, I wanted to create a CNN with multiple convolutional layers, each followed by a pooling operation to obtain an image with reduced height and width but with greater depth to obtain a significant detection of the most important features of each image just as defined by the guidelines. The **first convolutional layer** consists of **6 5x5x3** kernels. The third dimension set to 3 is obviously needed since RGB images have three channels to model the colors, while setting six small filters with width and height equal to 5 is a common effective choice for the first convolutional layer. I also set a **padding** initialized to 0 (of size 2 in this case, because the kernel width and height are set to 5) and stride equal to 1 so that the convolution is performed on all the pixels of the images. I made this choice because, given the limited dataset and the fact that the size of the images is quite small, I thought it was important not to leave out any pixels to avoid losing useful details. To introduce non-linearity, I used **Rectified Linear Units (ReLU)** activation function because it usually offers good performances and correctness. In addition, ReLU does not suffer significantly from saturation, and this facilitates network configuration and task completion. Now, to stabilize and speed up the model training, I added a **batch normalization** step. Finally, as previously mentioned, to obtain a reduction of the image, I inserted a **pooling** with size of 2x2 and stride of 2 which performs an average operation. I chose the average operation because, during the various tests I performed, the max operation led very quickly to overfitting. This may

mean that this type of task, with this type of architecture, benefits more from an average operation that enhances the entire area rather than from a max operation that gives greater importance to the most marked details. The second and third convolutional layers are analogous to the first, except that they use 16 and 32 kernels, respectively. Additionally, for the third layer, I set the kernel width and height to 3 and consequently set the padding size to 1, so that the kernel size remain significantly smaller than the output of the second convolutional layer and thus better capture the small characteristics. At this point, I noticed that with additional convolutional layers, or by further increasing the depth of the feature maps, I was getting significant overfitting. So, I decided to include the fully connected network architecture here. I therefore added the **flatten operation, two layers with 64 and 32 neurons respectively** and an **output layer with 5 neurons** since a classification into five classes must be performed. I also inserted a **20% dropout** after the first two layers to both reduce overfitting and not penalizing too much the training loss. The batch size is 64, the output activation function is obviously **softmax** and the loss function is **Cross Entropy Loss**. The Pytorch code snippet to create this architecture is the following (if Cross Entropy Loss is chosen, Pytorch will automatically use softmax after the output layer):

```
self.model = nn.Sequential(
    nn.Conv2d(3, 6, kernel_size = 5, stride = 1,
              padding = 2),
    nn.ReLU(),
    nn.BatchNorm2d(6),
    nn.AvgPool2d(2, stride = 2),

    nn.Conv2d(6, 16, kernel_size = 5, stride = 1,
              padding = 2),
    nn.ReLU(),
    nn.BatchNorm2d(16),
    nn.AvgPool2d(2, stride = 2),

    nn.Conv2d(16, 32, kernel_size = 3, stride = 1,
              padding = 1),
    nn.ReLU(),
    nn.BatchNorm2d(32),
    nn.AvgPool2d(2, stride = 2),

    nn.Flatten(),
    nn.Linear(32 * 12 * 12, 64),
    nn.ReLU(),
    nn.Dropout(0.2),

    nn.Linear(64, 32),
```

```

nn.ReLU(),
nn.Dropout(0.2),
nn.Linear(32, 5)
)

```

2.3 CNN1 with Adam optimizer

The first optimizer I chose is Adam, as it is generally known to bring good results. I set the **learning rate to 0.0001** because, from the various tests performed, I noticed that with $lr = 0.001$ the loss on the validation set, after a quite reduced number of epochs, began to oscillate. This may be a sign of the fact that, for this problem and with this architecture, it is a too high learning rate that leads to difficulties for the model to converge. Let us now present and describe the results obtained.

Class	Precision	Recall	F1-Score	Support
0	0.583	0.244	0.344	659
1	0.647	0.632	0.639	1005
2	0.650	0.730	0.687	988
3	0.549	0.759	0.637	1344
4	0.550	0.044	0.081	250
Macro Avg	0.596	0.482	0.478	4246
Weighted Avg	0.601	0.600	0.571	4246

Table 1: Performance metrics of the training phase.

Class	Precision	Recall	F1-Score	Support
0	0.669	0.284	0.399	341
1	0.671	0.519	0.585	495
2	0.673	0.713	0.693	512
3	0.514	0.823	0.633	656
4	1.000	0.025	0.049	119
Macro Avg	0.706	0.473	0.472	2123
Weighted Avg	0.641	0.594	0.566	2123

Table 2: Performance metrics of the validation phase.

Metric	Value
Accuracy (Training)	0.600
Accuracy (Validation)	0.594

Table 3: Accuracy of the training and validation phases.

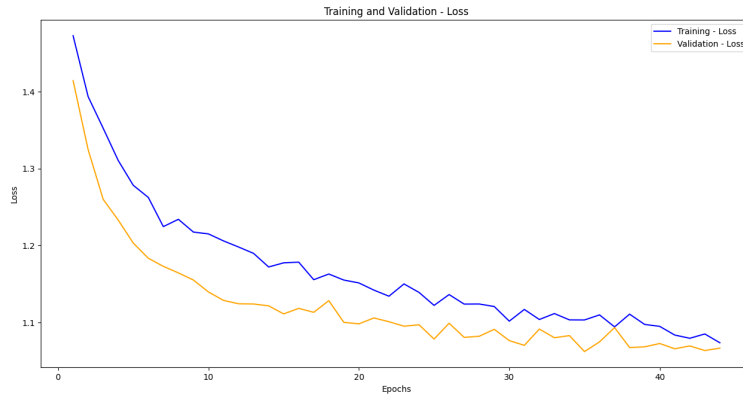


Figure 1: Training and validation loss.




Figure 2: Training and validation accuracy.



Figure 3: Training and validation F1-Score.

The training phase lasted **12 minutes** for a total of **44 epochs** as I noticed that, with a larger number of epochs, the training and validation losses, along with the other metrics, did not improve significantly and, on the contrary, the model started to tend towards overfitting. Now, let us see the test results.

Class	Precision	Recall	F1-Score	Support
0	0.387	0.323	0.352	133
1	0.396	0.458	0.425	275
2	0.517	0.692	0.592	406
3	0.836	0.778	0.806	1896
4	0.000	0.000	0.000	39
Macro Avg	0.427	0.450	0.435	2749
Weighted Avg	0.711	0.701	0.703	2749

Table 4: Performance metrics of the test phase.

Metric	Value
Accuracy (Test)	0.701

Table 5: Accuracy of the test phase.

Testing loss is **0.8798**.

Analyzing the results of the testing phase, we can notice, as anticipated, that the performances are not optimal. However, considering the use of CNN and the limitations of the training and test sets, we can conclude that, for this classification problem with five classes, an accuracy of 0.701 is not bad at all. Looking

at the other metrics, we can see how the quality of the results is proportional to the number of samples associated with the respective classes. In particular, class 3, the one with the most samples, shows good values of all the metrics: Precision, Recall and F1-Score. On the contrary, class 4, the one with only 39 samples, shows the worst performances and unfortunately the model was not able to correctly classify any image in this class. However, if we look at the performances of the validation phase, we can notice that the model was able to make correct predictions also for this class. This difference may be due to the very few samples of the test set. In particular, a Precision = 1 shows that the model correctly classified all the true positives it was able to find, but the low recall shows difficulty in identifying most of them. Following these results, it is logical that the macro average of the performances is significantly lower than the weighted average since the latter takes into account the number of samples for each class. The weighted average also shows significant improvement on the test set, demonstrating that the model can generalize quite well. Finally, the charts show how the loss correctly decreases following the gradient descent and how all the other metrics correctly encrease.

Let us now analyze the covariance matrix of the test phase.




Figure 4: Covariance Matrix.

The image shows that the results obtained are perfectly in line with what has just been said. An index of sufficient performance can be seen by noting that, except for classes 0 and 4 which are the least represented, the diagonal of the matrix shows that most of the predictions have been performed correctly. As expected, most of the classification errors are associated with the most represented classes, namely 2 and 3.

2.4 CNN1 with SGD and momentum

Let us now see the performance of this same architecture, with the same number of epochs, using SGD as the optimizer. To make the comparison, I decided to set the same learning rate as in the previous case, but adding a **momentum of 0.99** because SGD is generally more prone to getting stuck in local minima compared to Adam.

Class	Precision	Recall	F1-Score	Support
0	0.537	0.167	0.255	659
1	0.607	0.607	0.607	1005
2	0.627	0.725	0.672	988
3	0.534	0.751	0.624	1344
4	1.000	0.004	0.008	250
Macro Avg	0.661	0.451	0.433	4246
Weighted Avg	0.601	0.576	0.538	4246

Table 6: Training Results

Class	Precision	Recall	F1-Score	Support
0	0.630	0.284	0.392	341
1	0.697	0.432	0.534	495
2	0.682	0.746	0.713	512
3	0.502	0.843	0.629	656
4	0.000	0.000	0.000	119
Macro Avg	0.502	0.461	0.453	2123
Weighted Avg	0.583	0.587	0.554	2123

Table 7: Validation Results

Class	Precision	Recall	F1-Score	Support
0	0.423	0.308	0.357	133
1	0.430	0.378	0.402	275
2	0.493	0.727	0.588	406
3	0.828	0.792	0.810	1896
4	0.000	0.000	0.000	39
Macro Avg	0.435	0.441	0.431	2749
Weighted Avg	0.708	0.706	0.703	2749

Table 8: Test Results




Figure 5: Training and validation loss.




Figure 6: Training and validation accuracy.



Figure 7: Training and validation F1-Score.




Figure 8: Training and validation F1-Score.

The training phase took **12** minutes and the **loss is 0.9073**.

As we can see, the results obtained are very similar to the previous case. The most significant differences are the slightly higher loss and the fact that, although SGD managed to correctly predict more positive samples for classes 2 and 3, it obtained lower performances for the other classes. This could be due to the fact that SGD during the training phase learned even better the classification in these two classes at the expense of the correctness towards the others. As in the previous case, considering everything, the results obtained are still satisfactory.

3 Second approach

3.1 CNN architecture design

During the various tests I performed, I noticed that adding more convolutional layers resulted in a performance reduction, probably because the architecture

was starting to be too complex for the problem at hand. Therefore, for the second approach I decided to remove a convolutional layer. Additionally, I removed dropouts and added a L2 regularization with a parameter $\lambda = 0.01$. Since this architecture is "simpler" than the previous one, I decided to continue the training phase for **50 epochs**.

The Pytorch code snippet to create this architecture is the following:

```
self.model = nn.Sequential(
    nn.Conv2d(3, 6, kernel_size = 5, stride = 1,
              padding = 2),
    nn.ReLU(),
    nn.BatchNorm2d(6),
    nn.AvgPool2d(2, stride = 2),

    nn.Conv2d(6, 16, kernel_size = 3, stride = 1,
              padding = 1),
    nn.ReLU(),
    nn.BatchNorm2d(16),
    nn.AvgPool2d(2, stride = 2),

    nn.Flatten(),
    nn.Linear(16 * 24 * 24, 64),
    nn.ReLU(),

    nn.Linear(64, 32),
    nn.ReLU(),
    nn.Linear(32, 5)
)
```

3.2 CNN2 with Adam optimizer

Let's analyze the performance of this architecture using **Adam with lr = 0.001** as optimizer. Training phase took **8 minutes**.

Class	Precision	Recall	F1-Score	Support
0	0.632	0.287	0.395	659
1	0.666	0.614	0.639	1005
2	0.662	0.727	0.693	988
3	0.549	0.776	0.643	1344
4	0.706	0.096	0.169	250
Macro Avg	0.643	0.500	0.508	4246
Weighted Avg	0.625	0.610	0.587	4246

Table 9: Training Results

Class	Precision	Recall	F1-Score	Support
0	0.635	0.352	0.453	341
1	0.651	0.584	0.616	495
2	0.661	0.758	0.706	512
3	0.563	0.762	0.648	656
4	0.533	0.067	0.119	119
Macro Avg	0.609	0.505	0.508	2123
Weighted Avg	0.617	0.615	0.593	2123

Table 10: Validation Results

Class	Precision	Recall	F1-Score	Support
0	0.371	0.368	0.370	133
1	0.360	0.498	0.418	275
2	0.472	0.746	0.578	406
3	0.860	0.711	0.779	1896
4	0.080	0.051	0.062	39
Macro Avg	0.429	0.475	0.441	2749
Weighted Avg	0.718	0.669	0.683	2749

Table 11: Test Results

Dataset	Accuracy
Training	0.610
Validation	0.615
Test	0.669

Table 12: Accuracy for Training, Validation, and Test




Figure 9: Training and validation loss.




Figure 10: Training and validation accuracy.



Figure 11: Training and validation F1-Score.




Figure 12: Training and validation F1-Score.

Test loss is **0.9715**.

As we can observe, similar to the previous cases, the model achieved performances that, while not optimal, are still quite good. Charts show how the loss correctly decreases following the gradient descent and how all the other metrics correctly increase. It is interesting to notice that, differently from the previous models, here the training phase continued to improve for 50 epochs. This may be due to the fact that this network is shallower than the previous ones, so the training is able to last longer without starting overfitting. The most interesting aspect is that, although the metrics results are slightly lower than the previous ones, looking at the covariance matrix we can see that this model, except for class 3, is the one that correctly predicted the highest number of samples for each class. It also is the only one that has so far managed to correctly predict some samples for class 4. Most likely, this shallower network architecture made the model avoid focusing excessively on class 3 during the training phase, and this led to a better generalization ability towards less represented classes. The logical consequence is shown by the fact that the prediction errors concerning

class 3 also increased compared to the previous cases.

3.3 CNN2 with SGD and momentum

Finally, let us test this architecture with **SGD with lr = 0.0001** and **momentum = 0.99**. Training phase took **8 minutes** as before.

Class	Precision	Recall	F1-Score	Support
0	0.601	0.294	0.395	659
1	0.664	0.628	0.645	1005
2	0.657	0.699	0.678	988
3	0.549	0.774	0.642	1344
4	0.692	0.072	0.130	250
Macro Avg	0.633	0.493	0.498	4246
Weighted Avg	0.618	0.606	0.583	4246

Table 13: Training Results

Class	Precision	Recall	F1-Score	Support
0	0.641	0.287	0.397	341
1	0.581	0.667	0.621	495
2	0.627	0.828	0.714	512
3	0.584	0.646	0.614	656
4	0.000	0.000	0.000	119
Macro Avg	0.487	0.486	0.469	2123
Weighted Avg	0.570	0.601	0.570	2123

Table 14: Validation Results

Class	Precision	Recall	F1-Score	Support
0	0.310	0.301	0.305	133
1	0.338	0.596	0.432	275
2	0.417	0.791	0.546	406
3	0.864	0.622	0.724	1896
4	0.000	0.000	0.000	39
Macro Avg	0.386	0.462	0.401	2749
Weighted Avg	0.707	0.620	0.638	2749

Table 15: Test Results

Dataset	Accuracy
Training	0.606
Validation	0.601
Test	0.620

Table 16: Accuracy for Training, Validation, and Test




Figure 13: Training and validation loss.




Figure 14: Training and validation accuracy.



Figure 15: Training and validation F1-Score.




Figure 16: Training and validation F1-Score.

Test loss is **1.1157**.

As we can see, this last model is the one that performed the worst. However, the results achieved are still quite good: an accuracy of 0.620 and a weighted average of 0.638 with a CNN and limited and unbalanced training and test sets are not bad results. The graphs show that the loss is consistent with the gradient descent and that the other metrics show a steady growth. The loss tends not to improve significantly in the last epochs, which indicates that setting a higher early stopping can improve the results. The covariance matrix is consistent with what was just said. It is interesting to note that this model performed best of all the others exclusively for class 2. This could mean that, during training, the model focused too much on learning predictions for that class and this reduced the ability to generalize to the others.

4 Conclusion and future works

The results obtained show that, even in these cases, Adam obtained better results than SGD. Furthermore, it is important to notice that, for a problem where you have a limited and unbalanced dataset, a shallow CNN can perform better than a deep one because it will not focus excessively on the training data and will have a greater generalization capacity.

This task demonstrated how using a non-optimal method and having limited and unbalanced training and test sets can significantly lower performance and make it difficult to define a good architecture for the model. In particular, compared to the previous homework, in this case it was much more difficult to find the best architecture that I could identify as any change could worsen the performance or cause overfitting. It will be very interesting to solve this homework with Reinforcement Learning and compare the obtained performances.

5 Use of A.I.

For this homework, A.I. was only used to tabulate the values of the obtained metrics and to understand how to build a custom DataLoader class to apply only the data augmentation to the training set, as the official Pytorch documentation was not clear enough on this point.